

Topical Word Embedding

AAAI Conference on Artificial Intelligence, 2015

Yang Liu¹ Zhiyuan Liu¹ Tat-Seng Chuan² Maosong Sun^{1,3}

¹Department of Computer Science and Technology
State Key Lab on Intelligent Technology and Systems
National Lab for Information Science and Technology
Tsinghua University
Beijing 100084, China

²School of Computing, National University of Singapore
Singapore

³Jiangsu Collaborative Innovation Center for Language Competence
Jiangsu 221009, China

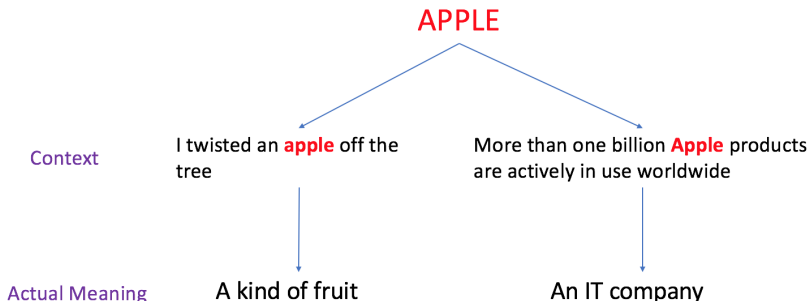


Plan

- 1 Introduction
 - Motivation
 - Related Work
- 2 Model
 - Topical Word Embeddings
- 3 Experiments
 - Contextual Word Similarity
 - Text Classification
 - Characteristics of TWE
- 4 Discussion



Motivation



- Most existing word embedding models :
 - Represent a word with a **single** vector
 - Indiscriminative for ubiquitous **homonymy** and **polysemy**



Motivation

- It is **problematic** that one word owns a **unique** vector for tasks like **Text Classification**
- Because many words have multiple senses
- To conceive a model that can **discriminate** word senses and generate multi-embeddings for each word



Related Work

- Skip-Gram model
- Multi-prototype vector model





Skip-Gram

Main idea :

- A well-known framework for learning word vectors
- Represent each word as a low-dimensional, dense vector via its context words
- If two words co-occur more frequently, then their word vectors are more similar, which is estimated by the cosine similarity of their word vectors



Linguistic Regularities in Skip-Gram




- Examples like – The famous "King - Man + Woman = Queen"



Skip-Gram Architecture

The quick brown fox jumps over the lazy dog.

- Aim to predict context words given a target word in a sliding window



Multi-prototype Vector Model

- **Cluster contexts** of a word into groups
- Generate a **distinct** prototype vector for each cluster




Multi-prototype Vector Model

1. Train single-prototype word vector

2. For **each word** with **one of its contexts**, average its neighboring word vectors as its **context feature vector**

The quick brown fox jumps over the lazy dog.




V_T

$$V_c = \frac{V_1 + V_2 + V_3 + V_4}{4}$$


Context Feature Vector

3. **Cluster** its context feature vectors into several prototypes(classes)

4. For each class, train one independent embedding for this word isolately




Disadvantages of Multi-prototype Vector Model



- Contexts of a word are divided into clusters with **no overlaps**
 - A word's several senses may **correlate** with each other – No clear semantic boundary
 - Like "**Play** the guitar" and "**Play** basketball"



Disadvantages of Multi-prototype Vector Model



- Generate multi-prototype vectors for each word in **isolation**
 - Ignore complicated correlations among words & their contexts



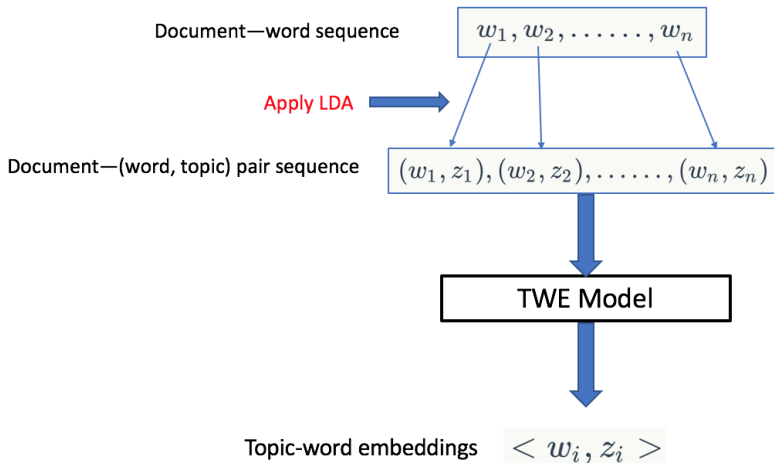
Topical Word Embeddings

- Basic idea is to allow each word to have **different embeddings under different topics**
- An example – for (word, topic) pairs :
 - (apple, food) indicates a fruit while (apple, IT) indicates an IT company



Topical Word Embeddings

Main process :



Skip-Gram

- Given a sequence $D = w_1, \dots, w_M$, maximize the average log probability (objective function)

$$\mathcal{L}(D) = \frac{1}{M} \sum_{i=1}^M \sum_{-k \leq c \leq k, c \neq 0} \log \Pr(w_{i+c} | w_i)$$

- Where the probability function is actually a softmax function as follows

$$\Pr(w_c | w_i) = \frac{e^{\vec{w}_c \cdot \vec{w}_i}}{\sum_{w_j \in W} e^{\vec{w}_c \cdot \vec{w}_j}}$$




Three TWE Models

- TWE-1
- TWE-2
- TWE-3



TWE-1



- Regard each topic as a **pseudo word**
- Learn topic embeddings and word embeddings **separately**
- **Concatenate** them to build a topical word embedding



TWE-1

- Regard each topic as a pseudo word
- Objective function :

$$\mathcal{L}(D) = \frac{1}{M} \sum_{i=1}^M \sum_{-k \leq c \leq k, c \neq 0} \log Pr(w_{i+c} | w_i) + \log Pr(w_{i+c} | z_i)$$

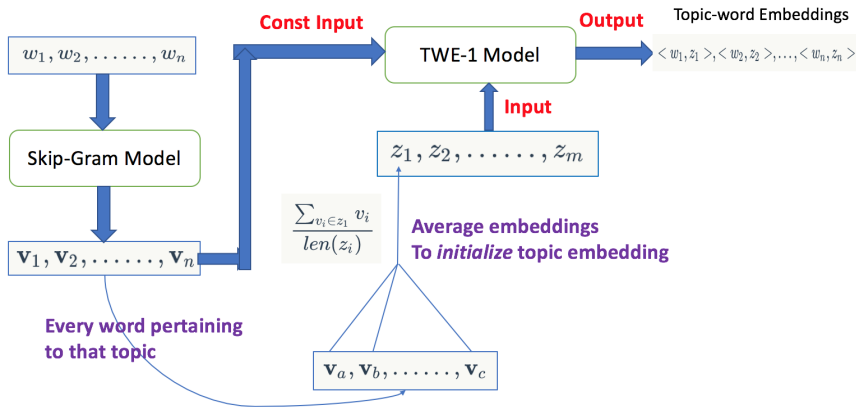
- Concatenate the embedding of w & z to get the topical word embedding :

$$\vec{w}^z = \vec{w} \oplus \vec{z}$$



TWE-1

Main Process :




TWE-1

Disadvantage

- Not consider the **immediate interaction** between a word and its assigned topic for learning



TWE-2



- Consider each word-pair $\langle w_i, z_i \rangle$ as a **pseudo word**
- Learn topical-word embeddings directly
- For word **apple** under topics **food** and **IT**, we consider $\langle \text{apple}, \text{food} \rangle$ is a **word** and $\langle \text{apple}, \text{IT} \rangle$ is **another word**



TWE-2

- Objective Function :

$$\mathcal{L}(D) = \frac{1}{M} \sum_{i=1}^M \sum_{-k \leq c \leq k, c \neq 0} \log \Pr(\langle w_{i+c}, z_{i+c} \rangle \mid \langle w_i, z_i \rangle)$$

- Where \Pr is also a softmax function

$$\Pr(w_c | w_i) = \frac{e^{\vec{w}_c^{z_c} \cdot \vec{w}_i^{z_i}}}{\sum_{\langle w_c, z_c \rangle \in \langle W, T \rangle} e^{\vec{w}_c^{z_c} \cdot \vec{w}_i^{z_i}}}$$



TWE-2

Main Process :

- Initialize each topic-word pair using Skip-Gram
- Learn TWE models



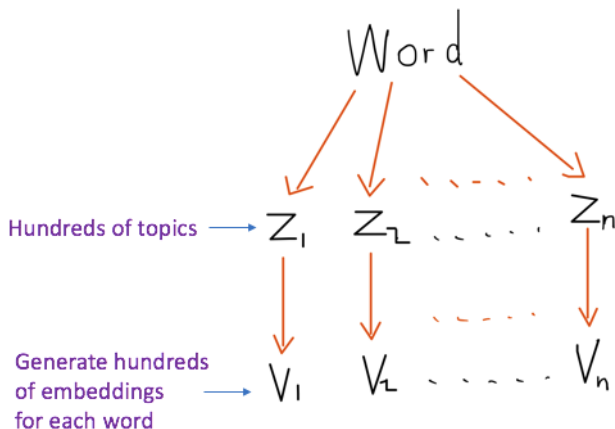
TWE-2

Disadvantage

- Although it considers the inner interaction of a word-topic pair,
- It suffers from the **sparsity issue**




TWE-2



- A word is actually **separated** into hundreds of parts



TWE-3



- Concatenate the vector \vec{w} and \vec{z} to build \vec{w}^z , with

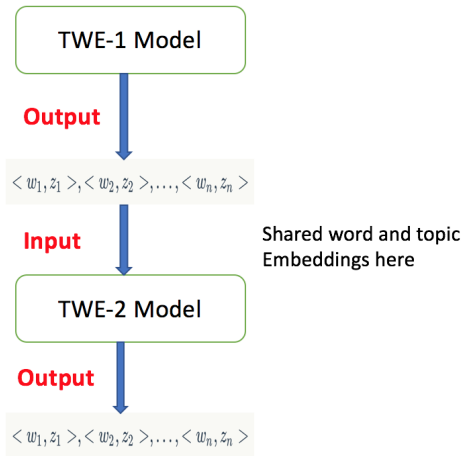
$$|\vec{w}^z| = |\vec{w}| + |\vec{z}|$$

- The objective of TWE-3 and probability function is identical to TWE-2, as shown above



TWE-3

Main Process :




TWE-3

TWE-2



TWE-3



ATTENTION

- In TWE-3, the parameter of each word embedding \vec{w} and topic embedding \vec{z} is **shared** over all word-topic $\langle w, z \rangle$
- Example : In TWE-2, two word-topic pair $\langle w, z \rangle$ and $\langle w, z' \rangle$ will have **distinct** parameters, while in TWE-3 they share the **same** word embedding


TWE-3

Disadvantage

- Although TWE-3 provide a trade-off between discrimination and sparsity
- It makes those words in the same topic **less discriminative**, because during the learning process, topic embeddings will influence the corresponding word embeddings



Contextual Word Similarity



Dataset :

- Our Dataset – SCWS
 - contain 2003 pairs of words and sentences containing these words
 - human labeled word similarity scores based on the word meaning in the context

Method :

- Compute **Spearman Correlation** between the human labeled result and the result from the model



Contextual Word Embedding

- For each word with its context c , infer the topic distribution by

$$Pr(z|w, c) \propto Pr(w|z)Pr(z|c)$$

- Further obtain the contextual word embedding as

$$\vec{w}^c = \sum_{z_i \in T} Pr(z_i|w, c) \vec{w}^{z_i}$$



Contextual Word Embedding

- AvgSimC method :

$$\begin{aligned} S(w_i, c_i, w_j, c_j) &= \vec{w}_i^{c_i} \cdot \vec{w}_j^{c_j} \\ &= \sum_{z \in T} \sum_{z' \in T} Pr(z|w_i, c_i) Pr(z'|w_j, c_j) S(\vec{w}^z, \vec{w}^{z'}) \end{aligned}$$

- MaxSimC method :

- $\vec{w}^c = \vec{w}^z, z = \arg \max_z Pr(z|w, c)$
- $S(w_i, c_i, w_j, c_j) = \vec{w}_i^z \cdot \vec{w}_j^{z'}$



Contextual Word Similarity

Compared Methods :

- Single-prototype model
 - C&W (Collobert and Weston 2008)
 - TFIDF & Pruned TFIDF
 - LDA-S
 - LDA-C
 - Skip-Gram
- Multi-prototype model
 - Pruned TFIDF-M
 - Huang's model (Huang et al. 2012)
 - Tian's model (Tian et al. 2014)



Contextual Word Similarity

Table 2: Spearman correlation $\rho \times 100$ of contextual word similarity on the SCWS data set.

Model	$\rho \times 100$	
C&W	57.0	
TFIDF	26.3	
Pruned TFIDF	62.5	
LDA-S	56.9	
LDA-C	50.4	
Skip-Gram	65.7	
	AvgSimC	MaxSimC
Pruned TFIDF-M	60.5	60.4
Tian	65.4	63.6
Huang	65.3	58.6
TWE-1	68.1	67.3
TWE-2	67.9	63.6
TWE-3	67.1	65.5



Analysis

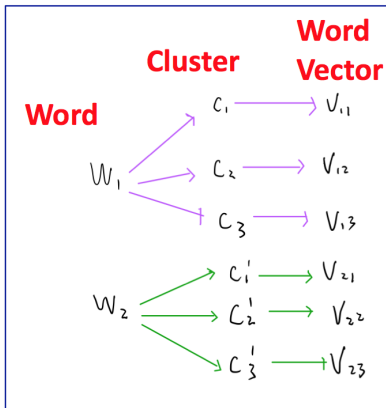
TWE performs better than other multi-prototype models

- More topics than clusters, more details of word semantics (Topics=400, cluster=8)
- Most multi-prototype models build multi-prototypes for each word **separately**, ignoring rich **interactions** between words as well as their contexts
- LDA provides a more **principle** way to select the most **appropriate** topical word embedding for a word under specific context




Analysis

Multi-prototype Model



TWE Model



Text Classification

Data Set :

- 20NewsGroup
 - 20, 000 documents from 20 different newsgroup

Method :

- Use Liblinear Classification
- Report macro-averaging precision, recall and F-measure for comparison



Document Embedding

- Represent the semantics of a document by aggregating over all topical word embeddings :

$$d = \sum_{w \in d} Pr(w|d) \vec{w}^z$$

- Where $Pr(w|d)$ can be weighted with TFIDF scores of words in d



Text Classification

Table 3: Evaluation results of multi-class text classification.

Model	Accuracy	Precision	Recall	F-measure
BOW	79.7	79.5	79.0	79.0
LDA	72.2	70.8	70.7	70.0
Skip-Gram	75.4	75.1	74.3	74.2
PV-DM	72.4	72.1	71.5	71.5
PV-DBOW	75.4	74.9	74.3	74.3
TWE-1	81.5	81.2	80.6	80.6
TWE-2	79.0	78.6	77.9	77.9
TWE-3	77.4	77.2	76.2	76.1



Analysis

TWE-1 achieves the best performance

- The reason may be the **independent assumption** in TWE-1
- The size of 20NewsGroup is to some extent small
 - TWE-2 and TWE-3 may achieve better performance given more data for training



Characteristics of TWE

Table 4: Nearest neighbor words by TWE-2 and Skip-Gram.

Words	Similar Words
bank bank#1 bank#2	citibank, investment, river insurance, stock, investor river, edge, coast
left left#1 left#2	right, leave, quit moved, arrived, leave right, bottom, hand
apple apple#1 apple#2	macintosh, ios, juice peach, juice, strawberry mac, ipod, android

- Find the most similar words of several example words in different topics



Discussion

- 1 Explore **non-parametric** topic models
 - In LDA, the topic number must be pre-defined
 - Use CV to find the topic number but are **time-consuming** and **impractical** for large-scale data
- 2 Is there any other **better embedding model** utilizing **topics**
 - The three TWE models somewhat have some disadvantages
 - And the last two models did not perform better as expected



Question ?



Thank you !

